

Track 6 | Session 3

如何藉由 AWS AI 和機器學習平台搭建 多功能的 AI 解決方案

Ivan Cheng (鄭志帆)

Manager, Solutions Architecture

Amazon Web Services

Centerpiece for digital transformation



Customer
experience



Business
operations



Decision
making



Innovation



Competitive
advantage

\$50B

By 2021, global spending on AI and cognitive technologies will exceed \$50 billion

Source: IDC, 2018

Our mission at AWS


Put machine learning (ML) in the
hands of every developer

Why AWS for ML?



Broadest and deepest set of AI and ML services

200+ new features and services launched in the last year



Accelerate your adoption of ML with Amazon SageMaker

Up to 70% cost reduction in data labeling

Up to 10x faster performance

Up to 75% lower inference cost

Reduced deployment time




Built on the most comprehensive cloud platform optimized for ML

AWS named as a leader in Gartner's Infrastructure as a Service (IaaS) Magic Quadrant for the 9th consecutive year

ML is happening in companies of every size and in every industry

10,000+ customers | More ML happens on AWS than anywhere else



Technology

Bringing AI into your digital transformation requires a new stack that makes it easier to put ML to work

The AWS ML stack

Broadest and most complete set of ML capabilities

AI services

Vision	Speech		Text			NEW! Search	Chatbots	Personalization	Forecasting	NEW! Fraud	NEW! Development	NEW! Contact centers
Amazon Rekognition	Amazon Polly	Amazon Transcribe +Medical NEW	Amazon Comprehend +Medical	Amazon Translate	Amazon Textract	Amazon Kendra	Amazon Lex	Amazon Personalize	Amazon Forecast	Amazon Fraud Detector	Amazon CodeGuru	Contact Lens <i>For Amazon Connect</i>

ML services

	Ground Truth		Augmented AI	ML Marketplace	Amazon SageMaker Studio IDE NEW!								Neo
Amazon SageMaker					Built-in algorithms	NEW! Notebooks	NEW! Experiments	Model training & tuning	NEW! Debugger	NEW! Autopilot	Model hosting	NEW! Model Monitor	











ML frameworks & infrastructure

				Deep learning AMLs & containers		GPUs & CPUs	Amazon Elastic Inference	AWS Inferentia	FPGA


The AWS ML stack

Broadest and most complete set of ML capabilities







AI services

Vision			Speech		Text			NEW! Search	Chatbots	Personalization	Forecasting	NEW! Fraud	NEW! Development	NEW! Contact centers
														
Amazon Rekognition			Amazon Polly		Amazon Transcribe <small>+Medical</small> NEW			Amazon Kendra	Amazon Lex	Amazon Personalize	Amazon Forecast	Amazon Fraud Detector	Amazon CodeGuru	Contact Lens <small>For Amazon Connect</small>

ML services

 Amazon SageMaker	Ground Truth	Augmented AI	ML Marketplace	Amazon SageMaker Studio IDE NEW!								Neo
				Built-in algorithms	NEW! Notebooks	NEW! Experiments	Model training & tuning	NEW! Debugger	NEW! Autopilot	Model hosting	NEW! Model Monitor	


ML frameworks & infrastructure

 TensorFlow			 GLUON		 Keras	Deep learning AMIs & containers	GPUs & CPUs	Amazon Elastic Inference	AWS Inferentia	FPGA
										

The ML workflow is iterative and complex


Bringing ML to all developers

Prepare




Collect and
prepare
training data

Build




Choose or
build an
ML algorithm


Train & tune



Set up and
manage
environments
for training




Train,
debug, and
tune models




Manage
training runs

Deploy & manage




Deploy
model in
production



Monitor
models




Validate
predictions



Scale manage
the production
environment


Amazon SageMaker helps you build, train, and deploy models

Prepare




Collect and
prepare
training data

Build




Choose or
build an
ML algorithm


Train & tune



Set up and
manage
environments
for training




Train,
debug, and
tune models




Manage
training runs

Deploy & manage




Deploy
model in
production



Monitor
models



Validate
predictions



Scale manage
the production
environment

Amazon SageMaker helps you build, train, and deploy models


Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```


Collect and prepare training data

Build




Choose or build an ML algorithm


Train & tune



Set up and manage environments for training




Train, debug, and tune models




Manage training runs


Deploy & manage




Deploy model in production



Monitor models



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm

Classification

- Linear Learner
- XGBoost
- KNN

Working with text

- BlazingText
- Supervised
- Unsupervised

Computer vision

- Image classification
- Object detection
- Semantic segmentation

Recommendation

- Factorization machines

Anomaly detection

- Random Cut Forests
- IP insights

Regression

- Linear Learner
- XGBoost
- KNN

Topic modeling

- LDA
- NTM

Forecasting

- DeepAR forecasting

Feature reduction

- PCA
- Object2Vec

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models




Choose or build an ML algorithm

Train & tune

One-click training



Set up and manage environments for training




Train, debug, and tune models




Manage training runs

Deploy & manage




Deploy model in production



Monitor models



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training




Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models



Manage training runs


Deploy & manage




Deploy model in production



Monitor models



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training



Set up and manage environments for training

Debugging and optimization




Train, debug, and tune models

Visually track and compare experiments




Manage training runs


Deploy & manage




Deploy model in production



Monitor models



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training



Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models

Visually track and compare experiments




Manage training runs

Deploy & manage

One-click deployment and automatic scaling



Deploy model in production



Monitor models

Validate predictions

Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training



Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models

Visually track and compare experiments



Manage training runs


Deploy & manage

One-click deployment and automatic scaling




Deploy model in production


Automatically spot concept drift



Monitor models



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training



Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models

Visually track and compare experiments



Manage training runs


Deploy & manage

One-click deployment and automatic scaling



Deploy model in production

Automatically spot concept drift




Monitor models

Add human review of predictions



Validate predictions



Scale manage the production environment

Web-based IDE for ML

Automatically build and train models

Amazon SageMaker helps you build, train, and deploy models

Prepare

Fully managed data processing jobs and data labeling workflows

```
101011010
010101010
000011110
```

Collect and prepare training data

Build

One-click notebooks, built-in algorithms and models



Choose or build an ML algorithm


Train & tune

One-click training



Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models

Visually track and compare experiments



Manage training runs


Deploy & manage

One-click deployment and automatic scaling



Deploy model in production

Automatically spot concept drift




Monitor models

Add human review of predictions



Validate predictions

Fully managed with automatic scaling for 75% less



Scale manage the production environment














Web-based IDE for ML

Automatically build and train models

The AWS ML stack

Broadest and most complete set of ML capabilities







AI services

Vision	Speech		Text			NEW! Search	Chatbots	Personalization	Forecasting	NEW! Fraud	NEW! Development	NEW! Contact centers
												
Amazon Rekognition	Amazon Polly	Amazon Transcribe <i>+Medical</i> NEW	Amazon Comprehend <i>+Medical</i>	Amazon Translate	Amazon Textract	Amazon Kendra	Amazon Lex	Amazon Personalize	Amazon Forecast	Amazon Fraud Detector	Amazon CodeGuru	Contact Lens <i>For Amazon Connect</i>

ML services

	Amazon SageMaker Studio IDE NEW!										Neo
Amazon SageMaker	Ground Truth	Augmented AI	ML Marketplace	Built-in algorithms	NEW! Notebooks	NEW! Experiments	Model training & tuning	NEW! Debugger	NEW! Autopilot	Model hosting	NEW! Model Monitor

ML frameworks & infrastructure

 TensorFlow		 GLUON	 Keras	Deep learning AMIs & containers		GPUs & CPUs	Amazon Elastic Inference	AWS Inferentia	FPGA
 PYTORCH									

New: Automate online fraud detection with Amazon Fraud Detector



Pre-built fraud
detection model
templates



Automatic
creation of
custom fraud
detection models



Models learn from
past attempts to
defraud Amazon



Amazon
SageMaker
integration



One interface to
review past
evaluations and
detection logic

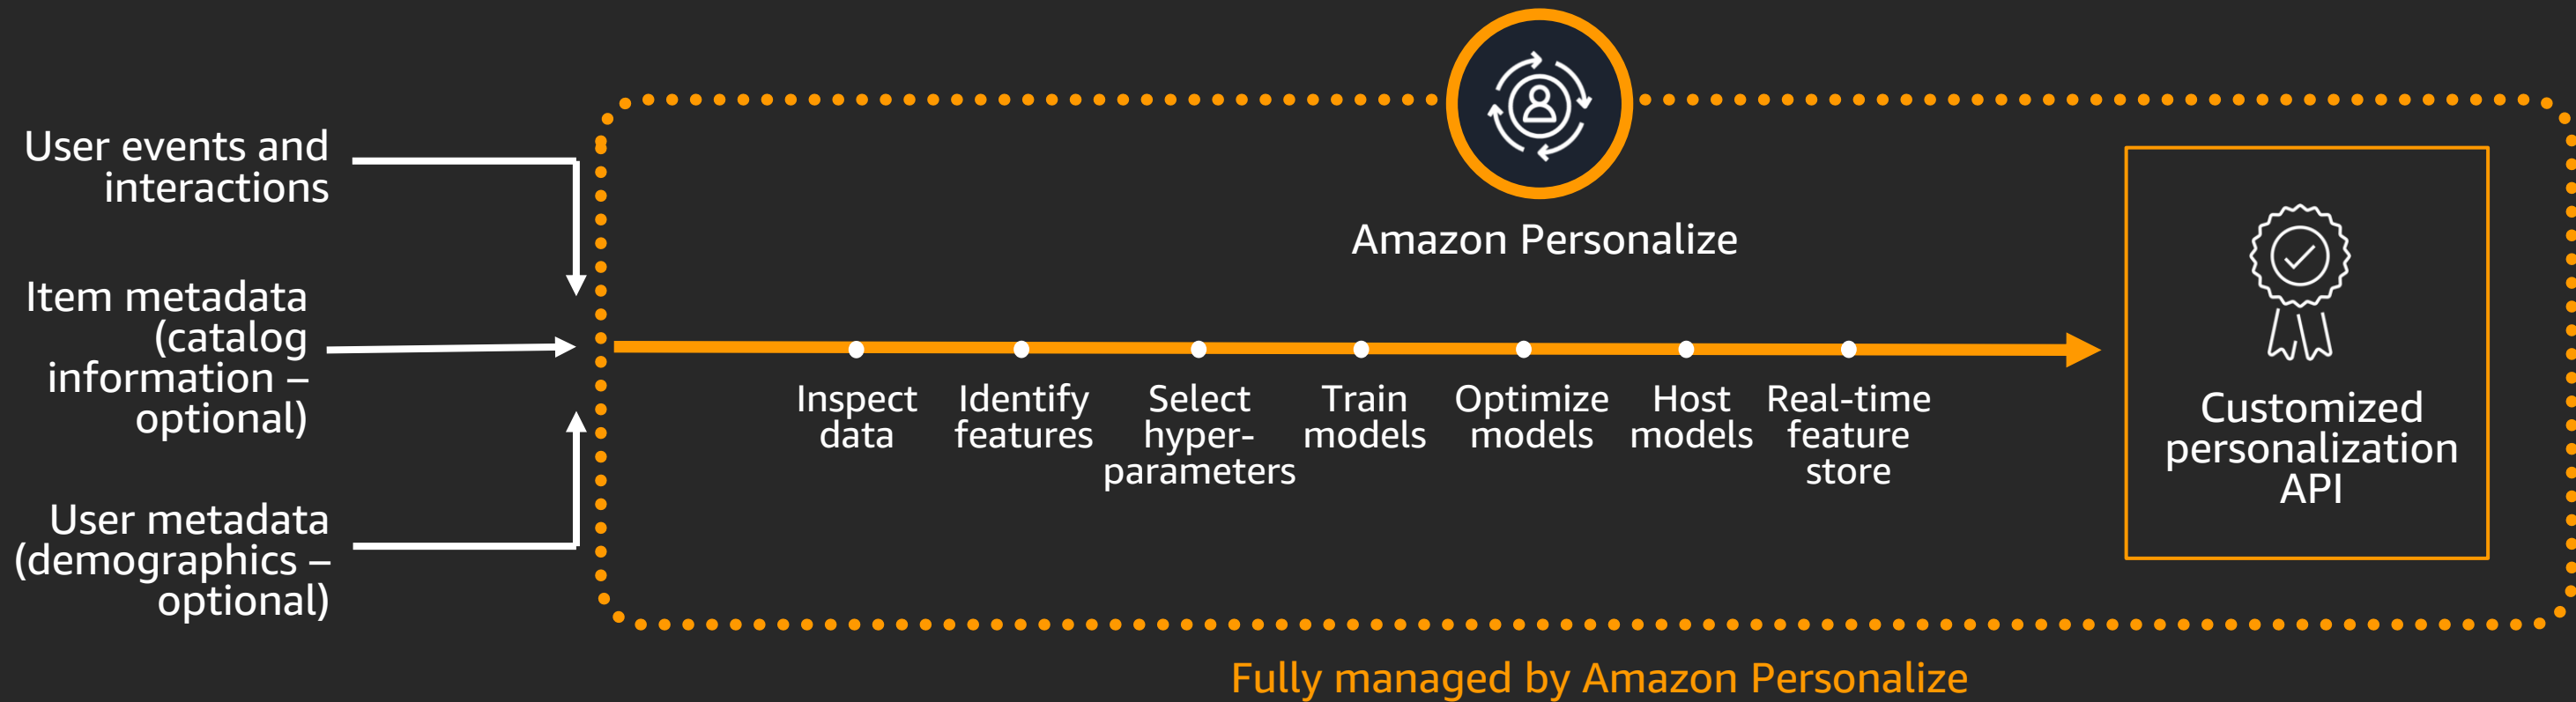
Amazon Personalize

Based on the technology that powers personalization at Amazon.com

- Fully managed service for generating personalized recommendations
- Uses the same ML technology that is being used at Amazon.com
- Generates highly relevant recommendations using deep learning techniques
- Does not require ML expertise to use
- Builds custom and private ML models using your own data

Amazon Personalize

Behind the scenes



Amazon Personalize

Improve customer experiences with personalization and recommendations



Deliver high-quality recommendations



Real-time



Deliver personalization in days, not months



Works with any product or content

Key features

Context-aware recommendations

Automated ML













Continuous learning to improve performance

Bring existing algorithms from Amazon SageMaker


The AWS ML stack

Broadest and most complete set of ML capabilities







AI services

Vision	Speech	Text			NEW! Search	Chatbots	Personalization	Forecasting	NEW! Fraud	NEW! Development	NEW! Contact centers
											
Amazon Rekognition	Amazon Polly	Amazon Transcribe +Medical NEW	Amazon Comprehend +Medical	Amazon Translate	Amazon Kendra	Amazon Lex	Amazon Personalize	Amazon Forecast	Amazon Fraud Detector	Amazon CodeGuru	Contact Lens <i>For Amazon Connect</i>

ML services

	Ground Truth	Augmented AI	ML Marketplace	Amazon SageMaker Studio IDE NEW!								Neo
				Built-in algorithms	NEW! Notebooks	NEW! Experiments	Model training & tuning	NEW! Debugger	NEW! Autopilot	Model hosting	NEW! Model Monitor	

ML frameworks & infrastructure

 TensorFlow	 mxnet	 GLUON	 Keras	Deep learning AMIs & containers	GPUs & CPUs	Amazon Elastic Inference	AWS Inferentia	FPGA
 PYTORCH								

Culture

Setting your organization up for success

Enabling the next ML developers

ML educational devices

AWS DeepLens
Deep learning




AWS DeepRacer
Reinforcement learning



AWS DeepComposer
Generative AI

Training and certification



AWS ML training
and certification

 UDACITY





Partnerships
with MOOCs

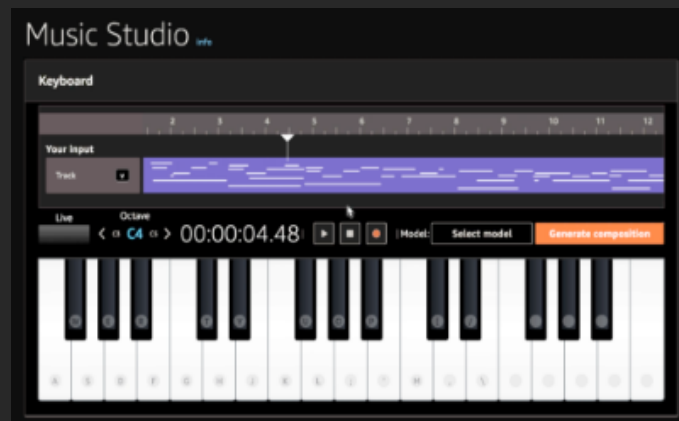
AWS DeepComposer

The world's first machine learning-enabled musical keyboard for developers



Creative meets generative

1



Input a melody by connecting the AWS DeepComposer keyboard

2




Choose from jazz, rock, pop, or classical, or build your own custom genre model in Amazon SageMaker

3



Publish your tracks to SoundCloud from the console; export MIDI files to your favorite DAW

AWS Machine Learning Solutions Lab



More than 100
deployed ML
engagements




Ideation
through to
production



Global
footprint





NEXT GEN STATS

powered by **aws**

Learn machine learning with AWS Training and Certification

Resources created by the experts at AWS to help you build and validate machine learning skills



Explore tailored machine learning (ML) paths for business decision makers, data platform engineers, data scientists, and developers



Learn at your convenience with 65+ free digital courses, or register for a live instructor-led class featuring hands-on labs and opportunities for practical application



Take the AWS Certified Machine Learning – Specialty exam to validate expertise in building, training, tuning, and deploying ML models

Visit the ML learning paths at <https://aws.training/ML>

Thank you!